

Level Up Your Django Performance: Identifying and Taming N+1 Queries



Mohammad Ahtasham ul Hassan

(Works at Arbisoft)



[mohammad-ahtasham](https://www.linkedin.com/in/mohammad-ahtasham)



Shafqat Farhan Ahmed

(Works at Arbisoft)



[shafqat-farhan-87677791](https://www.linkedin.com/in/shafqat-farhan-87677791)

Agenda

- Overview of ORMs with respect to relational databases
- Importance of performance in software applications
- What is N+1 problem
- Performance impact of N+1 problem
- Identifying N+1 queries (an example)
- Taming N+1 queries
- Considerations and key takeaways
- Questions?

Overview of ORMs with respect to relational databases

- Automatic query generation
- Django ORM operates on a lazy querying principle
- The database is only accessed when the query is actually evaluated
- N+1 query problem is particularly relevant to ORM and relational database queries
 - It can lead to performance issues

Importance of performance in software applications

- Why Performance matters?
 - User Experience
 - Timeouts
 - Random crashes/Memory overloads
 - Search Engine Rankings
 - Impact on Customer Retention
- Side effects of ORM on performance
 - Abstraction overhead
 - Increased Memory usage
 - N+1 problems

What is N+1 problem?

- When do N+1 queries occur?
 - Foreign Key Relationships
 - Reverse Foreign Key Lookups
 - Many-to-Many Relationships
 - Templates with ORM queries
 - Serializers with ORM queries
- Why doesn't the ORM handle itself?
 - **Performance:** Lazy loading helps to minimize the initial database query size
 - Complex relationships
 - **Flexibility:** Allows developers to optimize queries based on the specific use case

What is N+1 problem? (contd.)

```
class Conference(models.Model):  
    name = models.CharField(max_length=100)  
    date = models.DateField()  
  
class Talk(models.Model):  
    title = models.CharField(max_length=200)  
    duration = models.DurationField()  
    conference = models.ForeignKey(  
        Conference, related_name="talks", on_delete=models.CASCADE  
    )
```

What is N+1 problem? (contd.)

Code example:

```
✨ # Fetch all talks and their related conference data

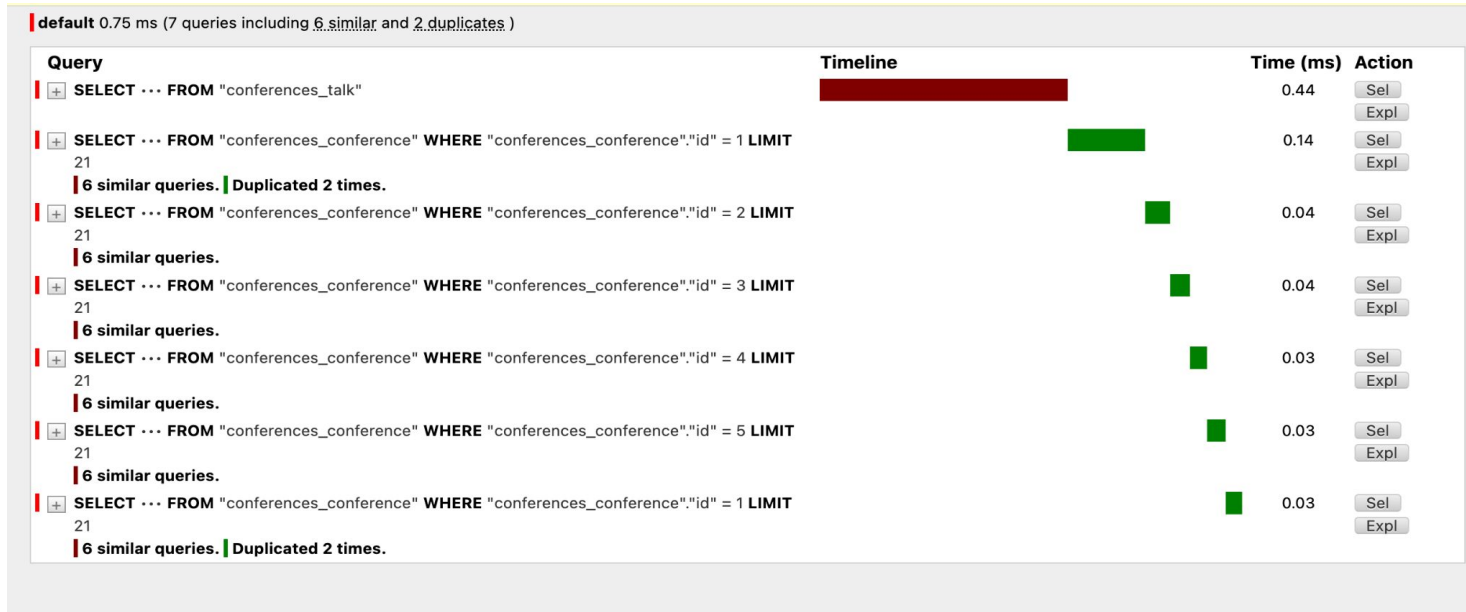
talks = Talk.objects.all()
conference_data = []
for talk in talks:
    conference_name = talk.conference.name
    conference_data.append(conference_name)
```

Performance impact of N+1 problem

- Increased database round trips for each query
- Degraded user experience
- Inefficient query execution
- Frequent API timeouts

Identifying N+1 queries (an example)

- Using Django Debug Toolbar
 - SQL query logging



Identifying N+1 queries (contd.)

- Code Reviews
- Testing
 - Use unit tests to test for executed queries

```
class MyTest(TestCase):  
    def test_query_count(self):  
        with self.assertNumQueries(1):  
            response = self.client.get('/books/')
```

- Other options?

Taming N+1 queries

- select_related()
 - For foreign key relations
 - Loads related objects in one query

```
talks = Talk.objects.select_related("conference").all()
```

- prefetch_related()
 - For many-to-many relations or reverse lookups
 - Performs an additional query and then join in Python

```
conferences = Conference.objects.prefetch_related('talks').all()
```

Taming N+1 queries (contd.)

default 0.61 ms (1 query)

Query	Timeline	Time (ms)	Action
<div><div></div><div><pre>SELECT "conferences_talk"."id", "conferences_talk"."title", "conferences_talk"."duration", "conferences_talk"."conference_id", "conferences_conference"."id", "conferences_conference"."name", "conferences_conference"."date" FROM "conferences_talk" INNER JOIN "conferences_conference" ON ("conferences_talk"."conference_id" = "conferences_conference"."id")</pre></div></div>		0.61	<div>SelExpl</div>
Connection: default			
<pre>/Users/mohammad.ahtasham/Desktop/djangocon/.venv/lib/python3.11/site-packages/django/contrib/staticfiles/handlers.py in __call__ (80) return self.application(environ, start_response)</pre>			
<pre>/Users/mohammad.ahtasham/Desktop/djangocon/conferences/views.py in conference_list_view (14) for talk in talks:</pre>			

Considerations and key takeaways

- The performance impact is minimal with small datasets
- Solving the N+1 problem may cause ripple effects in large and complex datasets
 - Prefetching large datasets can lead to increased memory usage
- Avoid prefetching related data if there's no specific use case (mindful data fetching)
- Django REST Framework (DRF) serializers fetching relational data can lead to N+1 query issues

Resources:

<https://docs.djangoproject.com/en/4.2/ref/models/querysets/#select-related>

<https://docs.djangoproject.com/en/4.2/ref/models/querysets/#prefetch-related>

Questions?